

HOME WORK 3 – Distributed Memory N-body with MPI

November 30, 2017

Introduction

This report is a basic summary of the analysis of N-body algorithm using OPenMPI for parallelization.

Code Modification

[Please click link to go to codes repo](#)

Benchmark Analysis

Table Key

PARTICLES: Number of Particles

PROCESSES: Number of Processes

TIME: Runtime in microseconds

SPEEDUP: $Sp = Ts / Tp$

Where;

Sp = Speed-up, Ts = Execution time in series Tp = Execution time in Parallel with P cores

EFFICIENCY: $Ep = Ts / pT$

Where; Ep = Efficiency, p = Number of Threads

Strong Scalability

Strong Scaling: Serial Benchmark

| PARTICLES | TIME |
|------------------|--------------|
| 1000 | 25.433723 |
| 2000 | 100.268770 |
| 4000 | 400.147935 |
| 8000 | 1604.038234 |
| 16000 | 6405.5837761 |

Strong Scaling: Parallel Benchmark for 1000 Particles

| PROCESSES | TIME | SPEEDUP | EFFICIENCY |
|-----------|----------|---------|------------|
| 2 | 13.2275 | 1.9228 | 0.9614 |
| 4 | 6.693263 | 3.7999 | 0.9500 |
| 8 | 3.358169 | 7.5737 | 1.0563 |
| 16 | 1.724769 | 14.746 | 0.9216 |
| 32 | 0.921790 | 27.592 | 0.8622 |
| 64 | 0.545337 | 46.639 | 0.7287 |
| 128 | 0.476963 | 53.324 | 0.4166 |
| 256 | 0.421122 | 60.395 | 0.2359 |

Strong Scaling: Parallel Benchmark for 2000 Particles

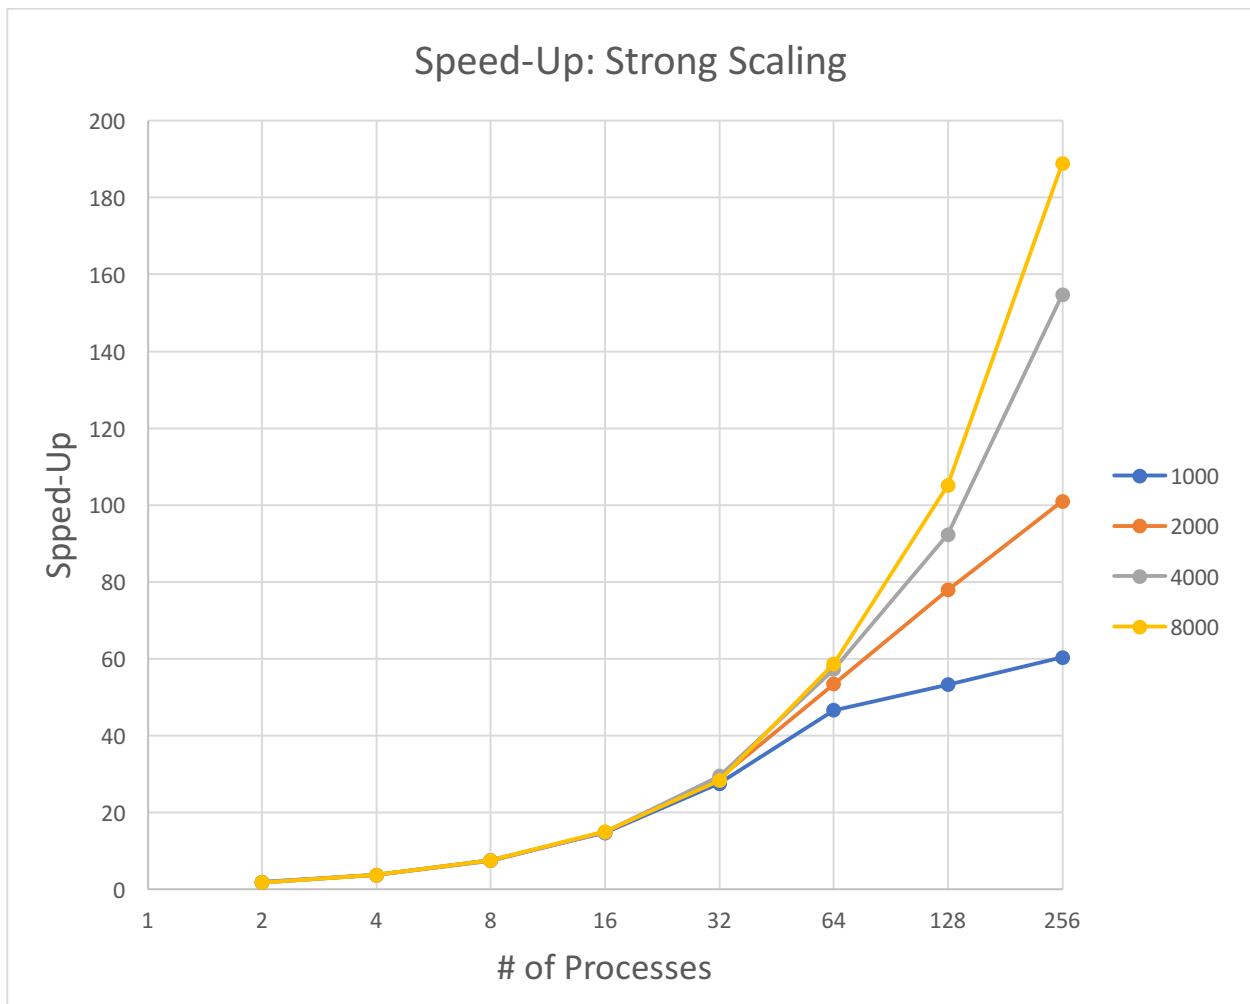
| PROCESSES | TIME | SPEEDUP | EFFICIENCY |
|-----------|-----------|---------|------------|
| 2 | 53.061749 | 1.8897 | 0.9448 |
| 4 | 26.636245 | 3.7644 | 1.0626 |
| 8 | 13.357500 | 7.5066 | 0.9383 |
| 16 | 6.729174 | 14.901 | 0.9313 |
| 32 | 3.468578 | 28.908 | 0.9034 |
| 64 | 1.874830 | 53.482 | 0.8339 |
| 128 | 1.286498 | 77.939 | 0.6089 |
| 256 | 0.992870 | 100.99 | 0.3945 |

Strong Scaling: Parallel Benchmark for 4000 Particles

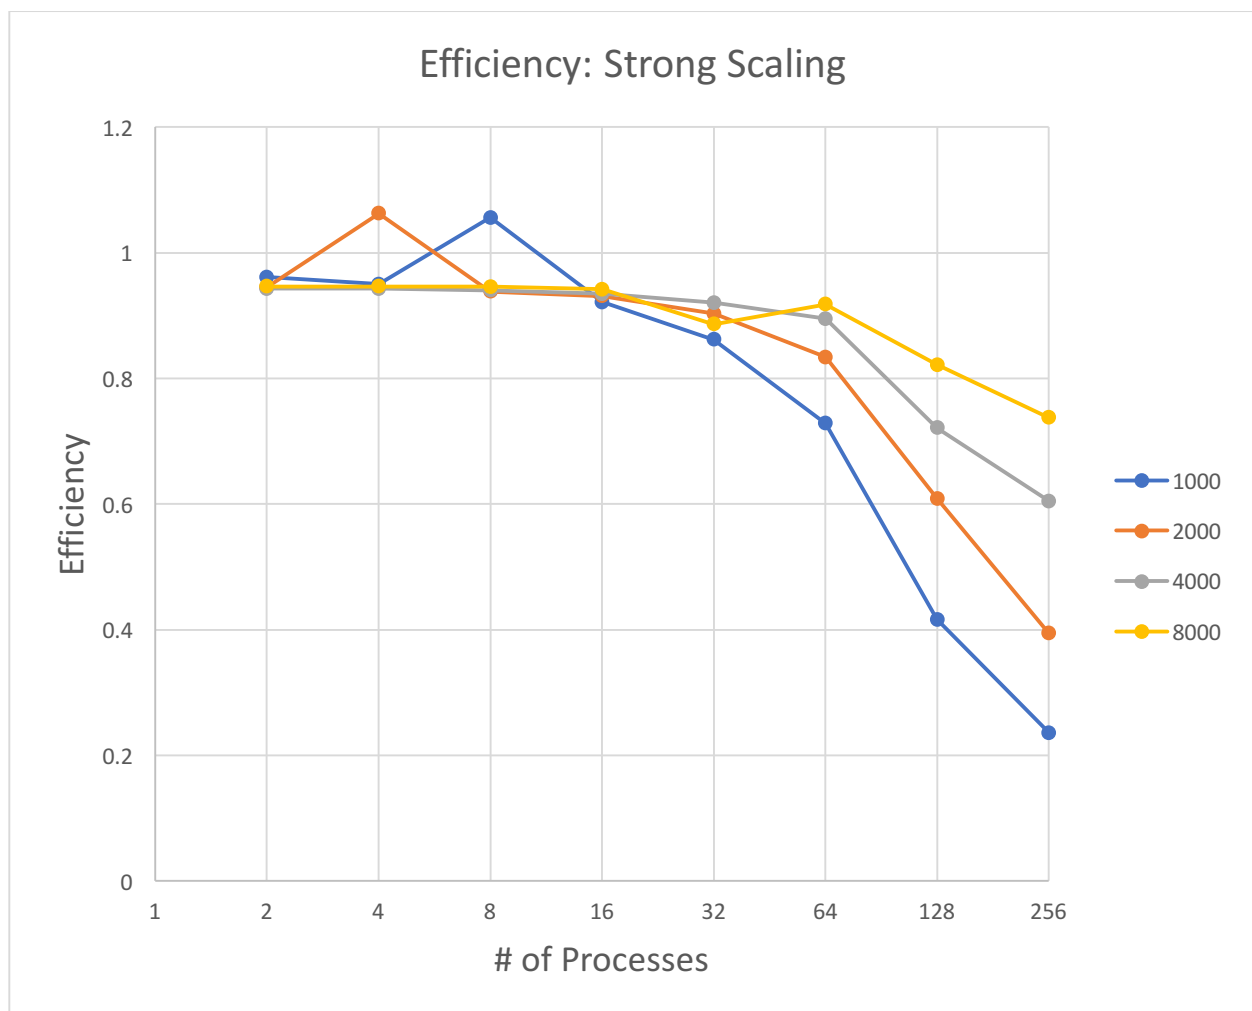
| PROCESSES | TIME | SPEEDUP | EFFICIENCY |
|-----------|------------|---------|------------|
| 2 | 212.146425 | 1.8862 | 0.9431 |
| 4 | 106.073760 | 3.7724 | 0.9431 |
| 8 | 53.191162 | 7.5228 | 0.9404 |
| 16 | 26.747779 | 14.960 | 0.9350 |
| 32 | 13.583481 | 29.458 | 0.9206 |
| 64 | 6.986296 | 57.276 | 0.8949 |
| 128 | 4.334127 | 92.325 | 0.7212 |
| 256 | 2.584722 | 154.81 | 0.6047 |

Strong scaling: Parallel Benchmark for 8000 Particles

| PROCESSES | TIME | SPEEDUP | EFFICIENCY |
|-----------|------------|---------|------------|
| 2 | 847.122028 | 1.8935 | 0.9468 |
| 4 | 423.612602 | 3.7866 | 0.9466 |
| 8 | 211.945913 | 7.5681 | 0.9460 |
| 16 | 106.412916 | 15.074 | 0.9421 |
| 32 | 56.556234 | 28.362 | 0.8863 |
| 64 | 27.303180 | 58.749 | 0.9180 |
| 128 | 15.250064 | 105.18 | 0.8217 |
| 256 | 8.490988 | 188.91 | 0.7379 |



In the graph above, for 1000 particles, we see an inclination from 2 processes through to just before 64 processes and a linear Speed-up from 64 processes to 256. We also see the same type of scaling for 2000, 4000, and 8000 particles as the Speed-up inclines from 2 processes through to 32, just before 128, and 128 processes respectively and a further Linear Speed-up to 256. We can infer that an increase in the number of process does not readily lead to a decline in the Speed-up even though the amount of work is reduced. Although the Speed-up is not continuously Super Linear all through, we still get a reasonable performance boost by increasing the number of processes.



From the graph above, we can see that for 1000 particles we see a Speed-up of about 95 percent using 2 processes and a very slight decline. Increasing the number of processes to 4 gives use a Linear performance boost and a further increment in the

number of processes causes an inverse performance, hence we see an almost perfect Linear decline. A further increment in the number of processes leads to a decline in performance. For 2000 particles, we notice about the same phenomenon with a Linear Speed-up at 2 processes from about 95 percent to about 106 percent. Increasing the number of the processes causes an almost Linear decline. A further increment in the number of processes causes a further decline. For Large enough work, like 4000 and 8000 particles, we don't see an increase, or decrease in performance with an increase in the number of processes until about 32 processes for 4000 particles and 16 processes for 8000 particles. Notice that at a further increment in the number of particles from 16 to 32 in 8000 particles gives a Linear Speed-up, until a further increment to 64 processes causes a reverse in performance, hence the decline.

Weak Scalability: Parallel Benchmark

| PARTICLES | PROCESSES | TIME | EFFICIENCY |
|------------------|------------------|-------------|-------------------|
| 200 | 1 | 1.067476 | 1 |
| 283 | 2 | 1.107529 | 0.96384 |
| 400 | 4 | 1.106645 | 0.96461 |
| 566 | 8 | 1.123555 | 0.95009 |
| 800 | 16 | 1.117051 | 0.95562 |
| 1131 | 32 | 1.164305 | 0.91684 |
| 1600 | 64 | 1.219269 | 0.87550 |
| 2263 | 128 | 1.594581 | 0.66944 |
| 3200 | 256 | 1.920175 | 0.55593 |

Weak Scaling: Parallel Benchmark


| PARTICLES | PROCESSES | TIME | EFFICIENCY |
|------------------|------------------|-------------|-------------------|
| 1000 | 1 | 25.495007 | 1 |
| 1414 | 2 | 26.611413 | 0.95805 |
| 2000 | 4 | 26.640297 | 0.95701 |
| 2828 | 8 | 26.680346 | 0.95557 |
| 4000 | 16 | 26.729546 | 0.95381 |
| 5657 | 32 | 26.914289 | 0.94740 |
| 8000 | 64 | 28.910461 | 0.88186 |
| 11314 | 128 | 30.709361 | 0.83020 |
| 16000 | 256 | 32.048332 | 0.79552 |

Weak Scaling: Parallel Benchmark

| PARTICLES | PROCESSES | TIME | EFFICIENCY |
|------------------|------------------|-------------|-------------------|
| 2000 | 1 | 102.036297 | 1 |
| 2828 | 2 | 106.170586 | 0.96106 |
| 4000 | 4 | 106.133483 | 0.96140 |
| 5657 | 8 | 106.399173 | 0.95900 |
| 8000 | 16 | 106.291287 | 0.96000 |
| 11314 | 32 | 106.933127 | 0.95421 |
| 16000 | 64 | 116.802235 | 0.87358 |
| 22627 | 128 | 121.197669 | 0.84190 |
| 31999 | 256 | 123.561676 | 0.82579 |

Weak Scaling: Parallel Benchmark

| PROCESSES | PARTICLES=200 | PARTICLES=1000 | PARTICLES=2000 |
|------------------|----------------------|-----------------------|-----------------------|
| 1 | 1.032528 | 25.975124 | 100.26877 |
| 2 | 0.560662 | 13.2275 | 53.061749 |
| 4 | 0.317195 | 6.693263 | 26.636245 |
| 8 | 0.169107 | 3.358169 | 13.357500 |
| 16 | 0.113186 | 1.724769 | 6.729174 |
| 32 | 0.093042 | 0.921790 | 3.468578 |
| 64 | 0.113931 | 0.545337 | 1.874830 |
| 128 | 0.151649 | 0.476963 | 1.286498 |
| 256 | 0.175286 | 0.421122 | 0.992870 |



The graph above is an illustration of the combined Weak Scaling Efficiency values of 200, 1000, and 2000 particles when parallelized with 1, 2, 4, 8, 16, 32, 64, 128, and 256 processes respectively.

As seen in the graph above, for all amount of work starting from 200, 1000 and 2000 particles respectively, Efficiency declines. A slight performance boost is gotten from increasing the number of processes in all cases and we observe a relative flat line (plato) through 4 to 8 and 16 processes. A further increment in the number of processes for a small amount of work like 200 particles causes the decline in performance, signifying there is not enough work to be done in parallel and the processes and spending too much time communicating than doing actual computing. The same is seen for larger amounts of work like 1000 and 2000 particles beginning at 32 processes through 256, shows a decline in performance with a relative increase in the number of processes

Conclusion

MPI scalability is relatively better than OpenMP scalability especially in the cases Speed-up and Efficiency for Strong Scaling. However, OpenMP seems to scale better in terms of Weak Scaling as it relates to Efficiency.

Challenges Faced

1. Finding the global(net) results across all ranks by implementing MPI Reduce
2. Implementing MPI_Allgatherv

Resolution:

- I. I emailed Dr. Stone for assistance and he walked me through it.